



Where should test-time reward go? Proposer weights vs. analyzer context vs. executor weights

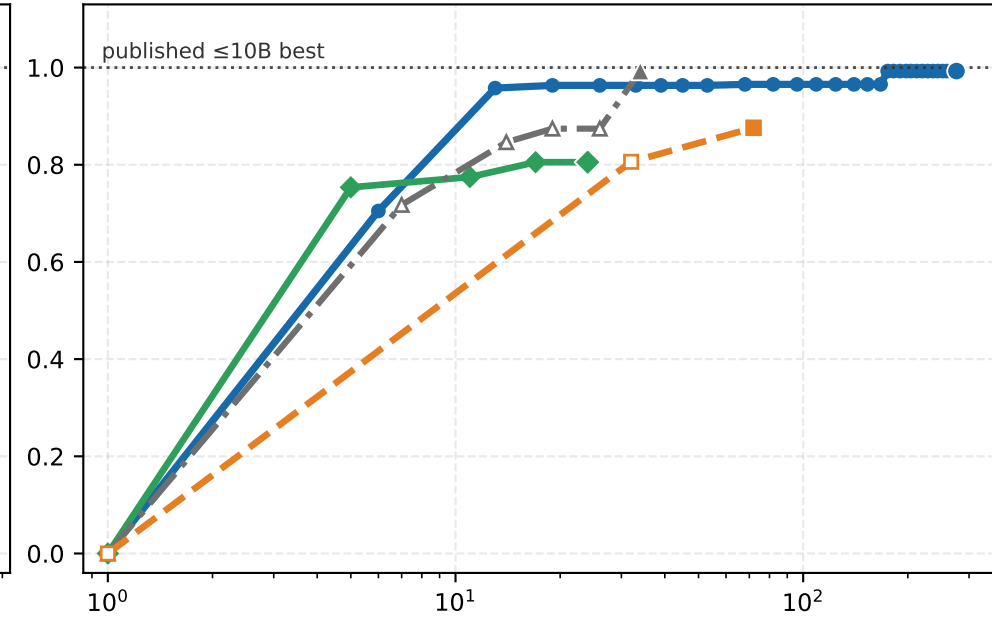
Erdős min-overlap †



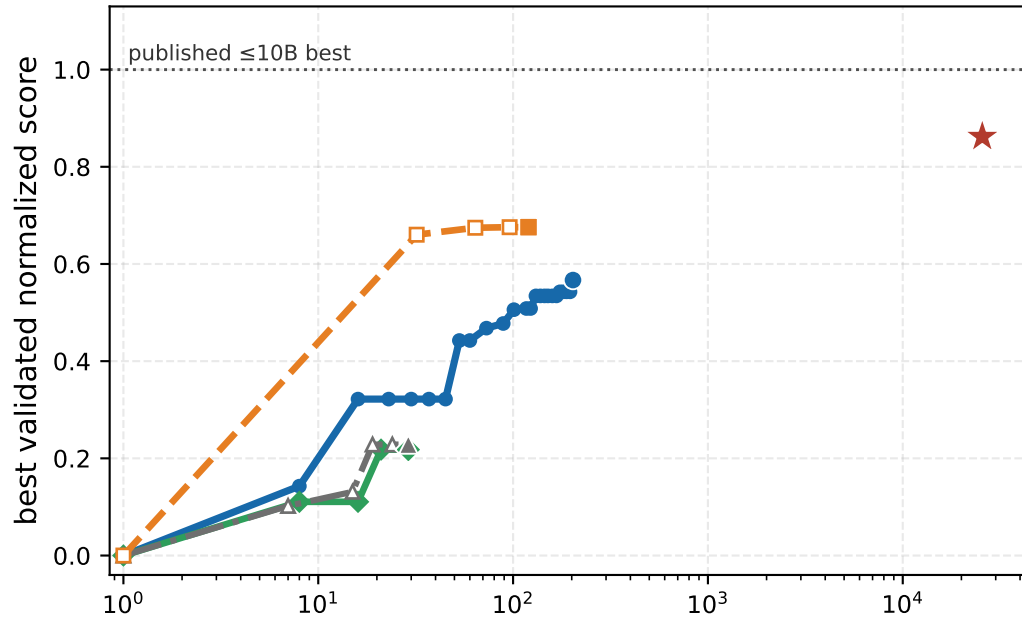
Circle packing (n=26)



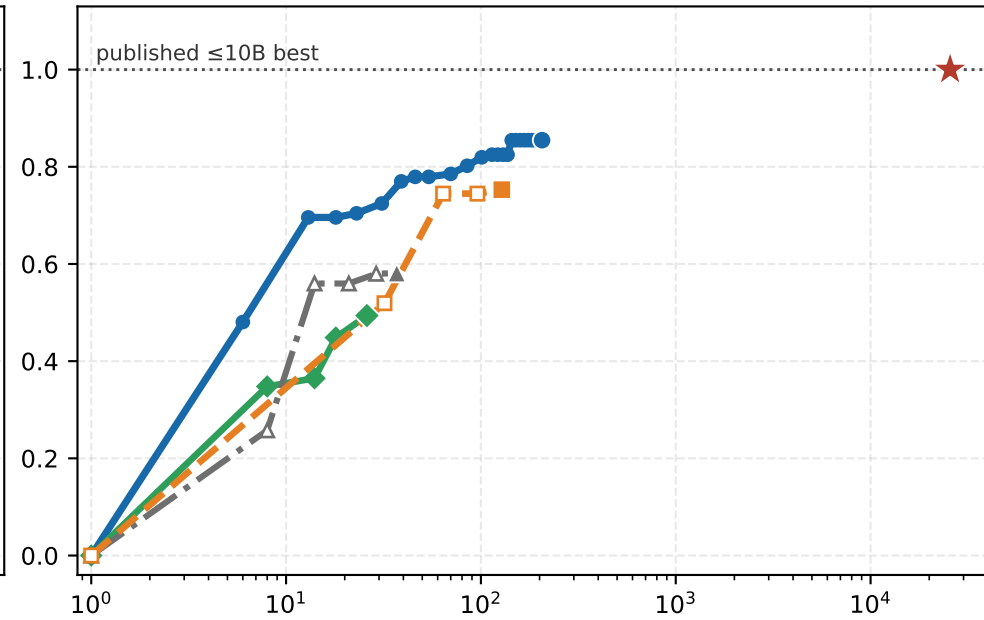
Hadamard max-det



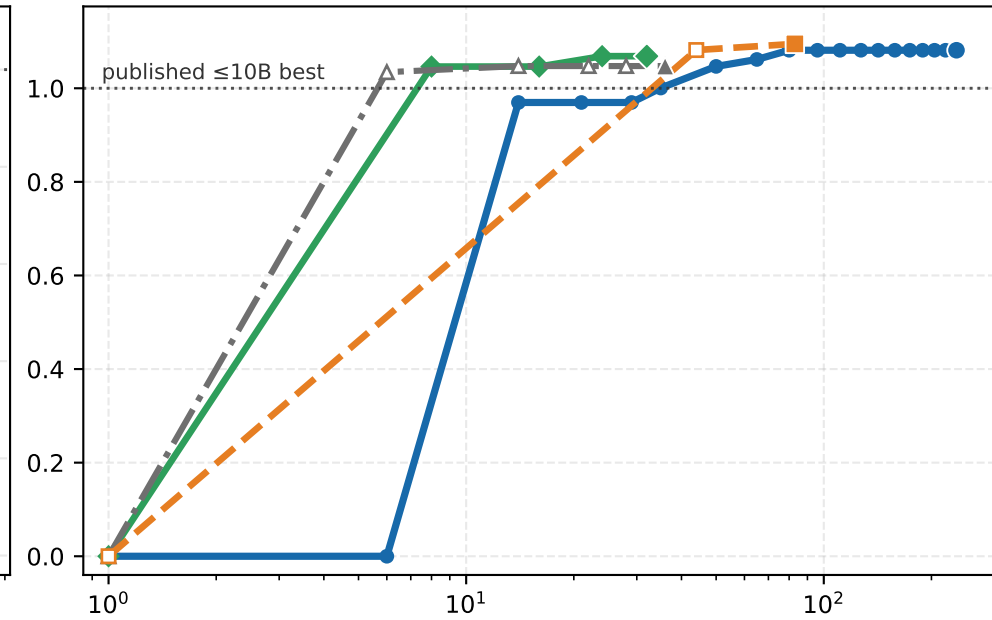
Autocorrelation I †



Autocorrelation II



AHC039



cumulative executor trajectories (log)

—●— Update proposer weights (ours, full campaign)
—◆— Update proposer weights (ours, matched arm)

cumulative executor trajectories (log)

—△— Update context (analyzer; frozen weights)
—□— Update executor weights (TTT-Discover-style)

★ TTT-Discover Qwen3-8B (published; 25.6k)

Green/gray/orange arms own independent ratchets; the blue full campaign is ceiling context and is excluded from common-budget metrics.

† Historical matched proposer/context arms used the wrong ratchet direction for Erdős and Autocorrelation I; shown for transparency, excluded from clean aggregate claims.

Solid endpoints are observed best-so-far values; no curve is extrapolated. Stars are published Qwen3-8B endpoints at 25.6k rollouts. Rollout efficiency excludes analyzer and training GPU cost; one campaign per matched arm, so no uncertainty estimate.