


Where should test-time reward go? Proposer weights vs. analyzer context vs. executor weights



Reward route	Inference per feedback batch	Local weight update
Proposer weights (ours)	K candidates: multi-call H1 proposal + executor trajectory	r64/a128 (116.4M trainable); 3 epochs; GBS 8; adapter continues, Adam resets
Analyzer context (ours ablation)	same H1 + executor path; 2 frozen-model analyst calls after cold round	none (proposer and executor weights frozen)
Executor weights (reference)	target K=8 executor trajectories; no H1/analyzer	r32/a64 (58.2M trainable); final-edit SFT; 1 epoch; GBS 4; adapter continues, Adam resets

Primary x = one shared fixed-H2 20-eval measurement at x=1 + actual launched executor trajectories; every fresh arm starts from task.initial_program.
Step curves credit a gain only after its batch completes; the vertical dashed line is each task's largest budget shared by all three routes.
Invalid H1 proposals, analyzer calls, evaluator caps/actual calls, optimizer boundaries, sandbox time, and allocated GPU-time proxies are reported separately in the audit.
* Blue inset values reproduce the proposer row in the main table (AHC039 curve: local 558,981 ledger; inset: separate 559,534 run without a local x). † Endpoint budgets differ; efficiency claims use common-budget score/AUC. Endpoints are observed, not extrapolated limits; unpaired single campaigns imply no confidence interval.
Reference note: the executor arm is a local budget-scaled TTT-Discover-style baseline, not an official reproduction; its final-edit replay and optimizer cadence differ from published TTT-Discover (full sampled trajectories, K=512).